

Introduction to R

[ECON 4753] — *University of Arkansas*

```
library(gapminder) # install.packages("gapminder")  
library(palmerpenguins) # install.packages("palmerpenguins")
```

1 — R as a Calculator

The first thing we will learn is how to use R as a calculator. You can use any of the math operators you want:

- + Addition
- - Subtraction
- * Multiplication
- / Division
- ^ Exponentiation

Let's experiment with some arithmetic expressions:

```
1 + 1
```

```
[1] 2
```

```
2 / 4
```

```
[1] 0.5
```

```
3^2
```

```
[1] 9
```

Order of operations (via PEMDAS) apply here too:

```
5 + 2 * 3
```

```
[1] 11
```

This does $2 * 3$ first and then adds 5 to get 11. If we want to do $5 + 2$ first, then we can wrap it in parenthesis (the P in PEMDAS):

```
(5 + 2) * 3
```

```
[1] 21
```

There are even some operators that you might not know that have to do with remainders:

```
# Get the remainder
```

```
13 %% 2
```

```
[1] 1
```

```
# Divide and round down to the nearest decimal
```

```
13 %/% 2
```

```
[1] 6
```

Exercise Compute the sample average of the following sample of baby weights (in lbs.):

(7.7, 8.2, 8.3, 7.6, 9.2, 7.4, 11.1)

```
## (7.7 + 8.2 + 8.3 + 7.6 + 9.2 + 7.4 + 11.1) / 7
```

```
## For variance, we can do:
```

```
## ((7.7 - 8.5)^2 + ..... ) / (7 - 1)
```

1.1. R as a functional programming language

R is based around functions. A function takes an input (or multiple inputs) and produces an output. There are many many functions in R, but first lets learn some calculator type functions. For example, if I want to take the square root, I can use the function `sqrt`. Here are some example of math functions:

The form of a function call is `function_name(arguments)`

1. The function name, `sqrt`, `abs`, `factorial`
2. Opening parenthesis (
3. The argument (in the future arguments)
4. Closing parenthesis)

For example `sqrt(16)` says to take the argument 16 and apply the function `sqrt` of it.

Exercise

1. Calculate the square root of 147
2. Try finding the natural log of 10, using `log()`
3. Practical usage: say the $Var(x) = 12$ and we have a sample size of 55. What is the standard deviation of the sample distribution of the sample mean?

1.2. Giving Things Names (i.e. Creating Variables)

Variables are immensely helpful in R. It lets you store values by giving them a name and then lets you access the variables later by name. I can assign variables using either `<-` or `=`.

Create variable `x` with value 5 and a variable `y` with value 20.

```
x <- 5
y <- 20
```

What is the sum of `x` and `y`?

```
x + y
```

```
[1] 25
```

Note the form of creating the variable:

1. The variable name, x and y
2. Assignment operator <- or =
3. The value we want to store.

The reason behind the left arrow is that the arrow points to variable name where we want to put the value into.

You can create a variable containing *text* by using " "

```
instructor_name <- "Kyle Butts"  
print(instructor_name)
```

```
[1] "Kyle Butts"
```

Exercise Use quotation marks to create a string and call it my_name.

```
my_name <- "Kyle Butts"  
print(my_name)
```

```
[1] "Kyle Butts"
```

1.3. Glueing together strings

Often time we might want to combine text from multiple sources and/or add data to our strings. We can use the paste/paste0 functions to combine together strings. paste0 will append the strings as written, while paste will automatically add a space between each thing it is concatenating. Dealer's choice for which you prefer

For example,

```
paste0("Kyle", "Butts")
```

```
[1] "KyleButts"
```

```
paste("Kyle", "Butts")
```

```
[1] "Kyle Butts"
```

What is cool about these functions is they take any number of arguments and append them together. For example, we can write a paste function that takes your height in inches and prints a human-readable string.

```
height_in_inches <- 70  
paste0("My height is ", height_in_inches %/% 12, "'", height_in_inches %% 12, "\"")
```

```
[1] "My height is 5'10\""
```

Note that numbers get automatically converted to a string. One subtle point you might have missed. If we start and end strings with double quotes, how can we include one in the text itself? Above, we did this with the escape key `\`. You probably won't need to do this, but it's worth mentioning nevertheless.

Exercise Construct a string that reports on the average baby weight in your sample. It would be nice to create a variable that stores `mean_baby_weight` to make the code nicer to read. These were the weights: (7.7, 8.2, 8.3, 7.6, 9.2, 7.4, 11.1)

2 — Vectors

So far, we have dealt with scalar numbers and texts. But, when working with data, we will observe many units and need a way to store all their values together. This is where the bread and butter of data-science comes in: Vectors.

Vectors are a list of elements like integers, numbers, or strings. This is really useful for storing data! You use `c` to create a vector (pneumonically, `c` stands for combine).

```
## Rebounds from 2023 NBA Season  
rebounds <- c(260, 114, 252, 310, 165, 236, 148, 336, 941, 127, 384, 278, 300, 6, 136,
```

It is kind of a pain in the neck to write all these out; and worse, prone to errors! Later, we will learn how to load data from a file, making this much easier.

You can access elements of a vector by using `[#]`, where `#` is the i -th element you want

```
rebounds[1]
```

```
[1] 260
```

```
rebounds[2]
```

```
[1] 114
```

If you want to access more than one element, we can subset using *a vector*! (how meta):

```
rebounds[c(1, 2)]
```

```
[1] 260 114
```

One special syntax is the `a:b` which generates $a, a + 1, \dots, b$. This makes it easy to grab the first 5 values:

```
rebounds[1:5]
```

```
[1] 260 114 252 310 165
```

Standard math operators work on vectors element by element:

```
rebounds[1:5] + 1
```

```
[1] 261 115 253 311 166
```

```
rebounds[1:5] / 12 # dozens of rebounds
```

```
[1] 21.66667  9.50000 21.00000 25.83333 13.75000
```

If we want to know how many elements are in a vector, we can use the `length` function:

```
length(rebounds)
```

```
[1] 386
```

2.1. Summarizing vectors

The natural next step is to start trying to summarize the data. There are a set of built-in functions that provide statistical summaries of the data.

```
## Mean, standard deviation, and variance
```

```
mean(rebounds)
```

```
[1] 250.4663
```

```
sd(rebounds)
```

```
[1] 232.6362
```

```
var(rebounds)
```

```
[1] 54119.62
```

```
## Extremes
```

```
max(rebounds)
```

```
[1] 1530
```

```
min(rebounds)
```

```
[1] 0
```

```
## chaining functions
```

```
sqrt(var(rebounds))
```

```
[1] 232.6362
```


Some functions will take extra **arguments** that give more instructions. For example, the `quantile` function returns information about percentiles of the distribution. You can add extra arguments, separated by commas. For example, `quantile`'s first argument is a vector and the second argument is what percentiles you want to find.

```
## Percentiles of the data  
quantile(rebounds, c(0.1, 0.5, 0.9))
```

```
10%    50%    90%  
19.0 205.5 575.5
```

We will talk in more details about plotting below, but for now we can use the `hist()` function to give us a view of the distribution of the vector

```
hist(rebounds)
```



2.1.1 Help menu

How did I know the order of arguments for `quantile`? Are there ways to customize the histogram created by `hist`? In general, if you have a question about a function, then you need to address

that function's **documentation**. To do so, click in to your console and type `?func_name` where `func_name` is the name of the function, e.g. `quantile`. At first, the information may be very overwhelming. The documentation in base R functions are very detailed, but are not great at introducing the functions. You should focus on the **Arguments** section and perhaps the **Examples** section, in my opinion. In particular, the order of the arguments is probably helpful in order of importance

Exercise Similar to `c`, the `seq` function creates a vector: a **sequence** of numbers.

1. Create a sequence of all multiples of four from 4 to 100. Look at `?seq` for help. Hint: The arguments you need here are `from`, `to`, and `by`. Store your vector in a variable
2. Find the 19th element of this sequence
3. What is the sum of the 10th and 11th element of this sequence.

2.1.2 NAs

In the real world, some times we will not have a value for a variable for an individual (e.g. people don't fill answer a survey question). In R, this is represented as an NA.

```
reviews <- c(5, NA, 4, 4, 3, 5, NA, 4, 5, 2)
```

What is the average (mean and median) review?

```
mean(reviews)
```

```
[1] NA
```

```
median(reviews)
```

```
[1] NA
```

By default, the statistical summary functions will all produce NA when they are present in the data. R wants you to *opt-in* to ignoring the missings. To do this, functions will take an extra **argument** called `na.rm`:

```
mean(reviews, na.rm = FALSE)
```

```
[1] NA
```

```
max(reviews, na.rm = FALSE)
```

```
[1] NA
```

A lot of this information will be presented to you by the summary function. For numeric vectors, summary produces the five-number summary, the mean, and the number of NAs (if any)

```
summary(reviews)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
2.00	3.75	4.00	4.00	5.00	5.00	2

2.1.3 Logical Vectors

So far, we have seen two kinds of vectors: numeric and character vectors. A third, common, vector is a **logical** vector:

```
ordered_takeout <- c(TRUE, TRUE, FALSE, TRUE, FALSE, TRUE, FALSE, FALSE)
```

Logical vectors can take only two values: TRUE and FALSE. They can be treated as numbers by using `as.numeric()` with TRUE becoming 1 and FALSE becoming 0.

```
as.numeric(ordered_takeout)
```

```
[1] 1 1 0 1 0 1 0 0
```

One common trick is to use `sum()` on a logical vector and it will return the number of TRUEs:

```
sum(ordered_takeout)
```

```
[1] 4
```

Logical vectors are often returned by other operations. For example, we can check whether elements of a vector equal some value with `==`. Other operators that produce a logical vector include `>`, `<`, `>=`, and `<=`.

```
## Find five-star reviews  
reviews == 5
```

```
[1] TRUE    NA FALSE FALSE FALSE  TRUE    NA FALSE  TRUE FALSE
```

```
## Find three-star or lower  
reviews <= 3
```

```
[1] FALSE    NA FALSE FALSE  TRUE FALSE    NA FALSE FALSE  TRUE
```

Logical vectors can be **negated** using the `!` operator:

```
## Find non five-star reviews  
!(reviews == 5)
```

```
[1] FALSE    NA  TRUE  TRUE  TRUE FALSE    NA  TRUE FALSE  TRUE
```

You can also use the following operators to supply multiple criteria:

- `&` And operator. Both vector 1 **and** vector 2 must be true for the observation
- `|` Or operator. **Either** vector 1 **or** vector 2 must be true for the observation

```
(reviews == 5) & (reviews == 4)
```

```
[1] FALSE    NA FALSE FALSE FALSE FALSE    NA FALSE FALSE FALSE
```

```
(reviews == 5) | (reviews == 4)
```

```
[1] TRUE NA TRUE TRUE FALSE TRUE NA TRUE TRUE FALSE
```

```
## Equivalent to
```

```
reviews >= 4
```

```
[1] TRUE NA TRUE TRUE FALSE TRUE NA TRUE TRUE FALSE
```

2.1.4 Subsetting of vectors by logical

We have already discussed one way of subsetting vectors via a vector of integers. This is called “subsetting by index”.

The other common way is by using a logical vector the same length of the vector you wish to subset. For example, let’s look at the reviews for takeout orders and dine-in orders separately

```
## Takeout orders
```

```
reviews[ordered_takeout]
```

```
[1] 5 NA 4 5 5 2
```

```
## Dine-in orders
```

```
reviews[!ordered_takeout]
```

```
[1] 4 3 NA 4
```

Exercise

1. What is the average review for take-out orders?
2. The `is.na()` function takes a vector as an argument and returns a logical vector that equals TRUE if the element is NA. Using this and the `!`, subset the reviews to only non-NA values. What is the average review? Compare this to `mean(reviews, na.rm = TRUE)`.

```
mean(reviews[!is.na(reviews)])
```

```
[1] 4
```

```
mean(reviews, na.rm = TRUE)
```

```
[1] 4
```

2.1.5 Vectorized operations

Often times we want to use multiple vectors for some calculation. Like single numbers, arithmetic can be done element-by-element with vectors. This means $+$, $-$, $*$, $/$ and $^$ all work on a vector.

```
x <- c(1, 2, 3)
```

```
y <- c(5, 5, 5)
```

```
x + y
```

```
[1] 6 7 8
```

```
x^2
```

```
[1] 1 4 9
```

There are two main features to remember: 1. The vectors you use should be of the same length (if they are not, some weird *recycling* rules occur that we will not discuss in this introduction). 2. Scalars are treated as a vector of the same length with that single number repeated for each element.

```
## Equivalent:
```

```
x + 2
```

```
[1] 3 4 5
```

```
x + rep(2, 3)
```

```
[1] 3 4 5
```

More, many functions are designed to be used on vectors element by element. All of the functions we used in the “Calculator” section do this:

```
exp(x)
```

```
[1] 2.718282 7.389056 20.085537
```

```
log(x)
```

```
[1] 0.0000000 0.6931472 1.0986123
```

```
sqrt(x^2)
```

```
[1] 1 2 3
```

Exercise

1. Try to guess the output of the following expression $2 \times x + y + 1$.

2.1.6 Sorting data

The final vector operation we will discuss is how to sort data; either ascending or descending in value. There are two ways to do this.

First, we can use the `sort()` function. The function takes a function and an optional decreasing argument. `decreasing` takes a logical TRUE/FALSE option. By default, it increases in values

```
sort(rebounds, decreasing = TRUE)
```

[1]	1530	1258	941	934	921	909	899	870	862	845	830	829	807	770	762
[16]	760	744	740	739	728	705	704	700	691	672	670	665	660	639	634
[31]	633	631	630	618	607	593	580	580	578	573	564	564	556	551	549
[46]	546	536	513	512	511	500	497	494	491	485	485	483	476	473	472
[61]	469	468	460	454	453	451	450	449	448	447	439	435	434	432	429
[76]	426	423	420	417	416	415	410	407	405	402	402	401	394	394	393
[91]	391	384	381	372	350	346	344	342	336	336	330	324	319	318	317
[106]	317	314	312	312	310	310	310	308	307	305	305	301	301	300	298
[121]	296	296	296	295	295	292	292	289	286	286	282	281	278	275	272
[136]	271	270	269	269	268	265	265	261	260	260	260	259	258	257	257
[151]	257	257	256	255	253	253	252	252	249	247	247	247	246	245	244
[166]	243	240	239	238	236	236	235	233	233	233	231	227	227	227	223
[181]	223	222	220	220	219	213	211	210	209	208	206	206	206	205	204
[196]	202	202	201	191	190	189	188	188	187	185	184	184	184	183	182
[211]	182	179	178	177	173	171	170	168	168	165	162	162	161	161	152
[226]	150	149	148	147	145	145	145	145	144	142	136	132	129	127	124
[241]	118	118	118	116	114	112	112	111	110	106	105	102	101	101	99
[256]	99	98	98	97	96	96	96	95	95	95	94	92	90	88	85
[271]	85	84	82	82	81	81	81	80	79	78	78	78	77	76	75
[286]	74	71	70	69	69	69	66	66	63	63	63	61	58	56	55
[301]	54	54	54	52	51	51	49	47	47	46	45	43	42	42	41
[316]	41	39	39	37	36	35	35	34	34	34	33	33	32	30	30
[331]	30	30	28	28	27	26	26	26	26	25	24	24	24	22	22
[346]	21	19	19	19	16	16	15	15	14	12	11	11	11	11	10
[361]	10	9	8	8	7	6	5	5	5	5	3	2	2	2	1
[376]	1	1	1	1	1	0	0	0	0	0	0				

Alternatively, we can use the `order()` function to get the *row indices* of the ordering:

```
order(rebounds, decreasing = TRUE)
```

```
[1] 298 379 9 313 113 215 177 244 347 252 25 295 139 310 141 85 216 111
[19] 98 58 258 371 250 209 88 245 181 130 255 317 183 82 272 71 373 158
[37] 65 127 41 219 217 271 263 265 64 197 273 30 356 180 96 97 238 101
[55] 23 39 213 49 235 79 349 227 228 372 234 195 240 332 249 150 108 330
[73] 33 138 194 266 60 19 327 294 315 87 163 377 129 189 67 204 285 282
[91] 320 11 344 364 292 62 44 144 8 291 229 63 287 233 179 283 184 174
[109] 262 4 148 311 205 106 169 385 286 319 13 103 32 190 303 261 369 175
[127] 176 93 61 123 376 43 12 159 91 152 333 36 381 230 334 359 193 1
[145] 131 384 83 268 140 210 290 363 323 275 109 348 3 232 208 92 342 378
[163] 353 40 56 221 383 34 361 6 366 37 17 214 270 134 90 107 203 55
[181] 362 38 120 212 297 116 105 115 117 218 162 198 370 188 316 104 318 53
[199] 331 124 72 99 325 70 146 74 136 165 339 126 178 248 100 321 277 299
[217] 149 112 307 5 31 42 22 202 66 118 284 7 289 16 119 243 374 122
[235] 18 15 225 375 10 300 132 206 309 324 2 27 171 246 242 280 128 157
[253] 76 358 46 155 151 167 257 95 259 306 69 145 276 367 207 352 355 156
[271] 211 52 47 237 26 322 368 173 48 80 182 354 147 326 386 251 281 89
[289] 51 73 137 196 312 29 94 201 114 68 125 350 241 260 301 54 164 279
[307] 187 302 338 191 267 170 143 314 81 154 45 264 351 77 20 199 57 75
[325] 226 86 236 185 172 224 253 304 329 360 254 121 192 239 305 220 110 142
[343] 160 35 341 365 288 293 345 28 222 256 380 308 84 21 135 186 278 161
[361] 168 382 153 269 78 14 59 200 274 335 50 223 296 346 24 102 231 247
[379] 328 336 133 166 337 340 343 357
```

The first element of the resulting vector is the index of the maximum number (1530)

```
order(rebounds, decreasing = TRUE)[1]
```

```
[1] 298
```



```
which.max(rebounds)
```

```
[1] 298
```

What this means is that we can use the result of `order` to subset the vector to get the sorted vector:

```
rebounds[order(rebounds, decreasing = TRUE)]
```

```
[1] 1530 1258 941 934 921 909 899 870 862 845 830 829 807 770 762
[16] 760 744 740 739 728 705 704 700 691 672 670 665 660 639 634
[31] 633 631 630 618 607 593 580 580 578 573 564 564 556 551 549
[46] 546 536 513 512 511 500 497 494 491 485 485 483 476 473 472
[61] 469 468 460 454 453 451 450 449 448 447 439 435 434 432 429
[76] 426 423 420 417 416 415 410 407 405 402 402 401 394 394 393
[91] 391 384 381 372 350 346 344 342 336 336 330 324 319 318 317
[106] 317 314 312 312 310 310 310 308 307 305 305 301 301 300 298
[121] 296 296 296 295 295 292 292 289 286 286 282 281 278 275 272
[136] 271 270 269 269 268 265 265 261 260 260 260 259 258 257 257
[151] 257 257 256 255 253 253 252 252 249 247 247 247 246 245 244
[166] 243 240 239 238 236 236 235 233 233 233 231 227 227 227 223
[181] 223 222 220 220 219 213 211 210 209 208 206 206 206 205 204
[196] 202 202 201 191 190 189 188 188 187 185 184 184 184 183 182
[211] 182 179 178 177 173 171 170 168 168 165 162 162 161 161 152
[226] 150 149 148 147 145 145 145 145 144 142 136 132 129 127 124
[241] 118 118 118 116 114 112 112 111 110 106 105 102 101 101 99
[256] 99 98 98 97 96 96 96 95 95 95 94 92 90 88 85
[271] 85 84 82 82 81 81 81 80 79 78 78 78 77 76 75
[286] 74 71 70 69 69 69 66 66 63 63 63 61 58 56 55
[301] 54 54 54 52 51 51 49 47 47 46 45 43 42 42 41
[316] 41 39 39 37 36 35 35 34 34 34 33 33 32 30 30
```

```
[331]  30  30  28  28  27  26  26  26  26  25  24  24  24  22  22
[346]  21  19  19  19  16  16  15  15  14  12  11  11  11  11  10
[361]  10   9   8   8   7   6   5   5   5   5   3   2   2   2   1
[376]   1   1   1   1   1   0   0   0   0   0   0
```

This might seem like a bit silly, but it will prove useful when we want to sort multiple vectors at the *same time* based on one of the vectors.

3 – Dataframes (or, a group of vectors)

Dataframes are a special object in R. A dataframe is simply a collection of **vectors** and looks like a typical excel spreadsheet. The columns of a dataframe are each **vectors** that contain variables and a row contains an **observation**. This is the coding equivalent of an excel spreadsheet. If you are using Positron, clicking the dataframe in the Variables tab or typing `View(df_name)` into the console will let you interactively scroll through the data.

First, we will load some data.frames that come with a **package** in R. We can do that using the `data` function. Let's load the penguins data set which contain a census conducted on multiple species of penguins on a set of islands.

```
data(penguins, package = "palmerpenguins")
```

We can use the `head()` function to view the first few rows. It prints out the first 6 rows of the dataset so you can see the variables. BTW, this function works on vectors too!

```
head(penguins)
```

```
# A tibble: 6 x 8
```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g
	<fct>	<fct>	<dbl>	<dbl>	<int>	<int>
1	Adelie	Torgersen	39.1	18.7	181	3750
2	Adelie	Torgersen	39.5	17.4	186	3800
3	Adelie	Torgersen	40.3	18	195	3250

```

4 Adelie Torgersen      NA      NA      NA      NA
5 Adelie Torgersen     36.7    19.3    193    3450
6 Adelie Torgersen     39.3    20.6    190    3650
# i 2 more variables: sex <fct>, year <int>

```

Another helpful function is `str()` which prints a similar format, but is a little easier to read, especially when there are a lot of variables in the dataset.

```
str(penguins)
```

```

tibble [344 x 8] (S3: tbl_df/tbl/data.frame)
 $ species      : Factor w/ 3 levels "Adelie","Chinstrap",...: 1 1 1 1 1 1 1 1 1 1 1 ...
 $ island       : Factor w/ 3 levels "Biscoe","Dream",...: 3 3 3 3 3 3 3 3 3 3 3 ...
 $ bill_length_mm : num [1:344] 39.1 39.5 40.3 NA 36.7 39.3 38.9 39.2 34.1 42 ...
 $ bill_depth_mm  : num [1:344] 18.7 17.4 18 NA 19.3 20.6 17.8 19.6 18.1 20.2 ...
 $ flipper_length_mm: int [1:344] 181 186 195 NA 193 190 181 195 193 190 ...
 $ body_mass_g    : int [1:344] 3750 3800 3250 NA 3450 3650 3625 4675 3475 4250 ...
 $ sex           : Factor w/ 2 levels "female","male": 2 1 1 NA 1 2 1 2 NA NA ...
 $ year          : int [1:344] 2007 2007 2007 2007 2007 2007 2007 2007 2007 2007 ...

```

Exercise

1. What constitutes a row in the penguins dataframe? What constitutes a column?

3.1. Accessing individual observations/variables

To access an individual variable, we can use the `$` operator. We use the `dataframe$variable` symbol to extract variable from the dataframe. For example, let's grab the species variable from penguins

Try grabbing the species and the sex variables us

```
penguins$species
```

```
[1] Adelie Adelie Adelie Adelie Adelie Adelie Adelie
```

[8]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[15]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[22]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[29]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[36]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[43]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[50]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[57]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[64]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[71]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[78]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[85]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[92]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[99]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[106]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[113]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[120]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[127]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[134]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[141]	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie	Adelie
[148]	Adelie	Adelie	Adelie	Adelie	Adelie	Gentoo	Gentoo
[155]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[162]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[169]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[176]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[183]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[190]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[197]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[204]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[211]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo

[218]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[225]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[232]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[239]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[246]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[253]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[260]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[267]	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo	Gentoo
[274]	Gentoo	Gentoo	Gentoo	Chinstrap	Chinstrap	Chinstrap	Chinstrap
[281]	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap
[288]	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap
[295]	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap
[302]	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap
[309]	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap
[316]	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap
[323]	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap
[330]	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap
[337]	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap	Chinstrap
[344]	Chinstrap						

Levels: Adelie Chinstrap Gentoo

penguins\$sex

[1]	male	female	female	<NA>	female	male	female	male	<NA>	<NA>
[11]	<NA>	<NA>	female	male	male	female	female	male	female	male
[21]	female	male	female	male	male	female	male	female	female	male
[31]	female	male	female	male	female	male	male	female	female	male
[41]	female	male	female	male	female	male	male	<NA>	female	male
[51]	female	male	female	male	female	male	female	male	female	male
[61]	female	male	female	male	female	male	female	male	female	male
[71]	female	male	female	male	female	male	female	male	female	male

```

[81] female male   female male   female male   male   female male   female
[91] female male   female male   female male   female male   female male
[101] female male   female male   female male   female male   female male
[111] female male   female male   female male   female male   female male
[121] female male   female male   female male   female male   female male
[131] female male   female male   female male   female male   female male
[141] female male   female male   female male   male   female female male
[151] female male   female male   female male   male   female female male
[161] female male   female male   female male   female male   female male
[171] female male   male   female female male   female male   <NA>   male
[181] female male   male   female female male   female male   female male
[191] female male   female male   female male   male   female female male
[201] female male   female male   female male   female male   female male
[211] female male   female male   female male   female male   <NA>   male
[221] female male   female male   male   female female male   female male
[231] female male   female male   female male   female male   female male
[241] female male   female male   female male   female male   male   female
[251] female male   female male   female male   <NA>   male   female male
[261] female male   female male   female male   female male   <NA>   male
[271] female <NA>   female male   female male   female male   male   female
[281] male   female female male   female male   female male   female male
[291] female male   male   female female male   female male   female male
[301] female male   female male   female male   female male   female male
[311] male   female female male   female male   male   female male   female
[321] female male   female male   male   female female male   female male
[331] female male   female male   male   female male   female female male
[341] female male   male   female
Levels: female male

```

The [,] operator wil let us subset rows and columns. Before comma = rows and After comma = columns

```
penguins[1:5, "species"]
```

```
# A tibble: 5 x 1
```

```
  species
```

```
  <fct>
```

```
1 Adelie
```

```
2 Adelie
```

```
3 Adelie
```

```
4 Adelie
```

```
5 Adelie
```

```
penguins[1:5,]
```

```
# A tibble: 5 x 8
```

```
  species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
```

```
  <fct>   <fct>         <dbl>         <dbl>             <int>         <int>
```

```
1 Adelie Torgersen      39.1           18.7             181           3750
```

```
2 Adelie Torgersen      39.5           17.4             186           3800
```

```
3 Adelie Torgersen      40.3            18             195           3250
```

```
4 Adelie Torgersen      NA              NA              NA              NA
```

```
5 Adelie Torgersen      36.7           19.3             193           3450
```

```
# i 2 more variables: sex <fct>, year <int>
```

```
penguins[, "species"]
```

```
# A tibble: 344 x 1
```

```
  species
```

```
  <fct>
```

```
1 Adelie
```

```
2 Adelie
```

```
3 Adelie
```

```
4 Adelie
5 Adelie
6 Adelie
7 Adelie
8 Adelie
9 Adelie
10 Adelie
# i 334 more rows
```

```
penguins[1:10, c("species", "island")]
```

```
# A tibble: 10 x 2
  species island
  <fct>   <fct>
1 Adelie Torgersen
2 Adelie Torgersen
3 Adelie Torgersen
4 Adelie Torgersen
5 Adelie Torgersen
6 Adelie Torgersen
7 Adelie Torgersen
8 Adelie Torgersen
9 Adelie Torgersen
10 Adelie Torgersen
```

You can pair these together, for example let's say I want the variable `island` for the first 6 observations:

```
penguins[1:6, "island"]
```

```
# A tibble: 6 x 1
  island
```



```
<fct>
1 Torgersen
2 Torgersen
3 Torgersen
4 Torgersen
5 Torgersen
6 Torgersen
```

```
penguins[1:6, ]$island
```

```
[1] Torgersen Torgersen Torgersen Torgersen Torgersen Torgersen
Levels: Biscoe Dream Torgersen
```

3.1.1 Exercise

1. Use the `unique()` function to find the unique values of the variable `species` in the `penguins` dataset.

```
unique(penguins$species)
```

```
[1] Adelie    Gentoo    Chinstrap
Levels: Adelie Chinstrap Gentoo
```

2. Use the `table()` function to find how many penguins there are of each species.

3.2. Loading data into R

In R, you can either load data from a website or from a computer. Usually data is found in a `.csv` file, but sometimes it will be in different forms that R can read.

```
# From a website
fandago <- read.csv("https://raw.githubusercontent.com/kylebutts/UARK_4753/refs/heads/main/data/fandago.csv")
head(fandago)
```

	FILM	RottenTomatoes	RottenTomatoes_User	Metacritic		
1	Avengers: Age of Ultron (2015)	74	86	66		
2	Cinderella (2015)	85	80	67		
3	Ant-Man (2015)	80	90	64		
4	Do You Believe? (2015)	18	84	22		
5	Hot Tub Time Machine 2 (2015)	14	28	29		
6	The Water Diviner (2015)	63	62	50		
	Metacritic_User	IMDB	Fandango_Stars	Fandango_Ratingvalue	RT_norm	RT_user_norm
1	7.1	7.8	5.0	4.5	3.70	4.3
2	7.5	7.1	5.0	4.5	4.25	4.0
3	8.1	7.8	5.0	4.5	4.00	4.5
4	4.7	5.4	5.0	4.5	0.90	4.2
5	3.4	5.1	3.5	3.0	0.70	1.4
6	6.8	7.2	4.5	4.0	3.15	3.1
	Metacritic_norm	Metacritic_user_norm	IMDB_norm	RT_norm_round		
1	3.30		3.55	3.90	3.5	
2	3.35		3.75	3.55	4.5	
3	3.20		4.05	3.90	4.0	
4	1.10		2.35	2.70	1.0	
5	1.45		1.70	2.55	0.5	
6	2.50		3.40	3.60	3.0	
	RT_user_norm_round	Metacritic_norm_round	Metacritic_user_norm_round			
1	4.5		3.5		3.5	
2	4.0		3.5		4.0	
3	4.5		3.0		4.0	
4	4.0		1.0		2.5	
5	1.5		1.5		1.5	
6	3.0		2.5		3.5	
	IMDB_norm_round	Metacritic_user_vote_count	IMDB_user_vote_count			
1	4.0		1330		271107	

2	3.5	249	65709
3	4.0	627	103660
4	2.5	31	3136
5	2.5	88	19560
6	3.5	34	39373

	Fandango_votes	Fandango_Difference
1	14846	0.5
2	12640	0.5
3	12055	0.5
4	1793	0.5
5	1021	0.5
6	397	0.5

However, most common is to download the data and put it in the folder where your .Rmd file is. To load data you will need to find the file location. Remember to set your working directory for this.

```
# penguins <- read.csv("penguins.csv")
```

4 – Distribution Functions

5 – Plotting

6 – Statistics